Capstonedatascience

September 13, 2020

[28]:
```
import pandas as pd
import pylab as pl
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
import matplotlib as mpl
```

[3]:
```
#let's import the example dataset from IBM
main_df=pd.read_csv ('https://s3.us.cloud-object-storage.appdomain.cloud/
 ↪cf-courses-data/CognitiveClass/DP0701EN/version-2/Data-Collisions.csv')
main_df.head ()
```

```
/home/jupyterlab/conda/envs/python/lib/python3.6/site-
packages/IPython/core/interactiveshell.py:3072: DtypeWarning: Columns (33) have
mixed types.Specify dtype option on import or set low_memory=False.
  interactivity=interactivity, compiler=compiler, result=result)
```

[3]:

	SEVERITYCODE	X	Y	OBJECTID	INCKEY	COLDETKEY	REPORTNO	\
0	2	-122.323148	47.703140	1	1307	1307	3502005	
1	1	-122.347294	47.647172	2	52200	52200	2607959	
2	1	-122.334540	47.607871	3	26700	26700	1482393	
3	1	-122.334803	47.604803	4	1144	1144	3503937	
4	2	-122.306426	47.545739	5	17700	17700	1807429	

	STATUS	ADDRTYPE	INTKEY	…	ROADCOND	LIGHTCOND	\
0	Matched	Intersection	37475.0	…	Wet	Daylight	
1	Matched	Block	NaN	…	Wet	Dark - Street Lights On	
2	Matched	Block	NaN	…	Dry	Daylight	
3	Matched	Block	NaN	…	Dry	Daylight	
4	Matched	Intersection	34387.0	…	Wet	Daylight	

	PEDROWNOTGRNT	SDOTCOLNUM	SPEEDING	ST_COLCODE	\
0	NaN	NaN	NaN	10	
1	NaN	6354039.0	NaN	11	
2	NaN	4323031.0	NaN	32	
3	NaN	NaN	NaN	23	
4	NaN	4028032.0	NaN	10	

```
                                        ST_COLDESC  SEGLANEKEY  \
0                                 Entering at angle           0
1  From same direction - both going straight - bo…           0
2                            One parked--one moving           0
3                From same direction - all others             0
4                                 Entering at angle           0

   CROSSWALKKEY  HITPARKEDCAR
0             0             N
1             0             N
2             0             N
3             0             N
4             0             N

[5 rows x 38 columns]
```

[4]:
```
#Now we will have a quick overview of the dataset we are dealing with
main_df.shape
```

[4]:
```
(194673, 38)
```

[5]:
```
#Looking at the columns we can determine what we need and what can be dropped
main_df.columns
```

[5]:
```
Index(['SEVERITYCODE', 'X', 'Y', 'OBJECTID', 'INCKEY', 'COLDETKEY', 'REPORTNO',
       'STATUS', 'ADDRTYPE', 'INTKEY', 'LOCATION', 'EXCEPTRSNCODE',
       'EXCEPTRSNDESC', 'SEVERITYCODE.1', 'SEVERITYDESC', 'COLLISIONTYPE',
       'PERSONCOUNT', 'PEDCOUNT', 'PEDCYLCOUNT', 'VEHCOUNT', 'INCDATE',
       'INCDTTM', 'JUNCTIONTYPE', 'SDOT_COLCODE', 'SDOT_COLDESC',
       'INATTENTIONIND', 'UNDERINFL', 'WEATHER', 'ROADCOND', 'LIGHTCOND',
       'PEDROWNOTGRNT', 'SDOTCOLNUM', 'SPEEDING', 'ST_COLCODE', 'ST_COLDESC',
       'SEGLANEKEY', 'CROSSWALKKEY', 'HITPARKEDCAR'],
      dtype='object')
```

0.1 Data Cleaning

Now that we have imported the data and had a quick verview of it we can begin the process of cleaning it and preparing it for the project

[6]:
```
#Here we drop all columns we do not need
main_df.drop(main_df.columns.difference(['SEVERITYCODE','SEVERITYDESC',␣
 ↪'ADDRTYPE', 'JUNCTIONTYPE', 'SDOT_COLDESC', 'WEATHER', 'LIGHTCOND',␣
 ↪'ROADCOND'])\
, axis=1, inplace=True)
main_df.head()
```

[6]:
```
   SEVERITYCODE      ADDRTYPE                    SEVERITYDESC  \
0             2  Intersection                Injury Collision
1             1         Block  Property Damage Only Collision
2             1         Block  Property Damage Only Collision
3             1         Block  Property Damage Only Collision
4             2  Intersection                Injury Collision

                             JUNCTIONTYPE  \
0   At Intersection (intersection related)
1  Mid-Block (not related to intersection)
2  Mid-Block (not related to intersection)
3  Mid-Block (not related to intersection)
4   At Intersection (intersection related)

                                       SDOT_COLDESC   WEATHER ROADCOND  \
0  MOTOR VEHICLE STRUCK MOTOR VEHICLE, FRONT END …  Overcast      Wet
1  MOTOR VEHICLE STRUCK MOTOR VEHICLE, LEFT SIDE …   Raining      Wet
2       MOTOR VEHICLE STRUCK MOTOR VEHICLE, REAR END  Overcast      Dry
3  MOTOR VEHICLE STRUCK MOTOR VEHICLE, FRONT END …     Clear      Dry
4  MOTOR VEHICLE STRUCK MOTOR VEHICLE, FRONT END …   Raining      Wet

                LIGHTCOND
0                Daylight
1  Dark - Street Lights On
2                Daylight
3                Daylight
4                Daylight
```

[7]:
```
#This provides a cleaner view of the columns we are going to use
main_df.head(0).transpose()
```

[7]:
```
Empty DataFrame
Columns: []
Index: [SEVERITYCODE, ADDRTYPE, SEVERITYDESC, JUNCTIONTYPE, SDOT_COLDESC,
WEATHER, ROADCOND, LIGHTCOND]
```

[8]:
```
#Now let's check for null values using boolean results
null_values=main_df.isnull()
null_values
```

[8]:
```
        SEVERITYCODE  ADDRTYPE  SEVERITYDESC  JUNCTIONTYPE  SDOT_COLDESC  \
0              False     False         False         False         False
1              False     False         False         False         False
2              False     False         False         False         False
3              False     False         False         False         False
4              False     False         False         False         False
…                  …         …             …             …             …
```

```
194668         False     False         False         False         False
194669         False     False         False         False         False
194670         False     False         False         False         False
194671         False     False         False         False         False
194672         False     False         False         False         False

        WEATHER  ROADCOND  LIGHTCOND
0         False     False      False
1         False     False      False
2         False     False      False
3         False     False      False
4         False     False      False
…           …         …          …
194668    False     False      False
194669    False     False      False
194670    False     False      False
194671    False     False      False
194672    False     False      False

[194673 rows x 8 columns]
```

[9]:
```
# We will check for null elements
for column in null_values.columns.values.tolist():
    print(column)
    print(null_values[column].value_counts().sort_values(ascending=True))
    print("")
```

```
SEVERITYCODE
False    194673
Name: SEVERITYCODE, dtype: int64

ADDRTYPE
True       1926
False    192747
Name: ADDRTYPE, dtype: int64

SEVERITYDESC
False    194673
Name: SEVERITYDESC, dtype: int64

JUNCTIONTYPE
True       6329
False    188344
Name: JUNCTIONTYPE, dtype: int64

SDOT_COLDESC
False    194673
```

```
Name: SDOT_COLDESC, dtype: int64

WEATHER
True       5081
False    189592
Name: WEATHER, dtype: int64

ROADCOND
True       5012
False    189661
Name: ROADCOND, dtype: int64

LIGHTCOND
True       5170
False    189503
Name: LIGHTCOND, dtype: int64
```

[10]:
```
#gives statistics for categorical variables
main_df.describe(include='O')
```

[10]:
```
        ADDRTYPE                    SEVERITYDESC  \
count     192747                          194673
unique         3                               2
top        Block  Property Damage Only Collision
freq      126926                          136485

                                    JUNCTIONTYPE  \
count                                     188344
unique                                         7
top     Mid-Block (not related to intersection)
freq                                       89800

                                                SDOT_COLDESC WEATHER ROADCOND  \
count                                                194673  189592   189661
unique                                                   39      11        9
top     MOTOR VEHICLE STRUCK MOTOR VEHICLE, FRONT END …   Clear      Dry
freq                                                  85209  111135   124510

        LIGHTCOND
count      189503
unique          9
top      Daylight
freq       116137
```

[11]:
```
main_df_with_nans=main_df.dropna()
main_df_with_nans.shape
```

[11]: (182914, 8)

[12]:
```
a=(1-(182954/194673))*100
print("%.2f" % a,"%")
```

```
6.02 %
```

[13]:
```
# With 6.02% rows with nans, we will drop these rows
main_df=main_df.dropna()
main_df.shape
```

[13]: (182914, 8)

[14]:
```
main_df.head()
```

[14]:

	SEVERITYCODE	ADDRTYPE	SEVERITYDESC	JUNCTIONTYPE	SDOT_COLDESC	WEATHER	ROADCOND	LIGHTCOND
0	2	Intersection	Injury Collision	At Intersection (intersection related)	MOTOR VEHICLE STRUCK MOTOR VEHICLE, FRONT END …	Overcast	Wet	Daylight
1	1	Block	Property Damage Only Collision	Mid-Block (not related to intersection)	MOTOR VEHICLE STRUCK MOTOR VEHICLE, LEFT SIDE …	Raining	Wet	Dark - Street Lights On
2	1	Block	Property Damage Only Collision	Mid-Block (not related to intersection)	MOTOR VEHICLE STRUCK MOTOR VEHICLE, REAR END	Overcast	Dry	Daylight
3	1	Block	Property Damage Only Collision	Mid-Block (not related to intersection)	MOTOR VEHICLE STRUCK MOTOR VEHICLE, FRONT END …	Clear	Dry	Daylight
4	2	Intersection	Injury Collision	At Intersection (intersection related)	MOTOR VEHICLE STRUCK MOTOR VEHICLE, FRONT END …	Raining	Wet	Daylight

[15]:
```
# we will reset index to correct rows numbers
main_df=main_df.reset_index(drop=True)
main_df.head()
```

[15]:
```
   SEVERITYCODE      ADDRTYPE                    SEVERITYDESC  \
0             2  Intersection                Injury Collision
1             1         Block  Property Damage Only Collision
2             1         Block  Property Damage Only Collision
3             1         Block  Property Damage Only Collision
4             2  Intersection                Injury Collision

                              JUNCTIONTYPE  \
0   At Intersection (intersection related)
1  Mid-Block (not related to intersection)
2  Mid-Block (not related to intersection)
3  Mid-Block (not related to intersection)
4   At Intersection (intersection related)

                                        SDOT_COLDESC   WEATHER ROADCOND  \
0  MOTOR VEHICLE STRUCK MOTOR VEHICLE, FRONT END …  Overcast      Wet
1  MOTOR VEHICLE STRUCK MOTOR VEHICLE, LEFT SIDE …   Raining      Wet
2       MOTOR VEHICLE STRUCK MOTOR VEHICLE, REAR END  Overcast      Dry
3  MOTOR VEHICLE STRUCK MOTOR VEHICLE, FRONT END …     Clear      Dry
4  MOTOR VEHICLE STRUCK MOTOR VEHICLE, FRONT END …   Raining      Wet

                LIGHTCOND
0                Daylight
1  Dark - Street Lights On
2                Daylight
3                Daylight
4                Daylight
```

[27]:
```
#Check numerical values of data
print ("SEVERITYCODE: \n", main_df ['SEVERITYCODE'].value_counts())
print("ADDRTYPE: \n", main_df['ADDRTYPE'].value_counts() )
print("LIGHTCOND: \n",main_df['LIGHTCOND'].value_counts())

print("\n WEATHER: \n",main_df['WEATHER'].value_counts())
print("\n JUNCTIONTYPE: \n",main_df['JUNCTIONTYPE'].value_counts())
print("\n SDOT_COLDESC: \n",main_df['SDOT_COLDESC'].value_counts())
print("\n ROADCOND: \n",main_df['ROADCOND'].value_counts())
```

```
SEVERITYCODE:
 1    126276
2     56638
Name: SEVERITYCODE, dtype: int64
ADDRTYPE:
 Block           119366
Intersection     63313
Alley              235
Name: ADDRTYPE, dtype: int64
```

```
LIGHTCOND:
 Daylight                    113850
Dark - Street Lights On      47550
Unknown                      10448
Dusk                          5772
Dawn                          2454
Dark - No Street Lights       1462
Dark - Street Lights Off      1157
Other                          210
Dark - Unknown Lighting         11
Name: LIGHTCOND, dtype: int64

 WEATHER:
 Clear                       109065
Raining                      32649
Overcast                     27189
Unknown                      11637
Snowing                        881
Other                          746
Fog/Smog/Smoke                 556
Sleet/Hail/Freezing Rain       112
Blowing Sand/Dirt               49
Severe Crosswind                25
Partly Cloudy                    5
Name: WEATHER, dtype: int64

 JUNCTIONTYPE:
 Mid-Block (not related to intersection)             86613
At Intersection (intersection related)               61221
Mid-Block (but intersection related)                 22341
Driveway Junction                                    10519
At Intersection (but not related to intersection)     2055
Ramp Junction                                          160
Unknown                                                  5
Name: JUNCTIONTYPE, dtype: int64

 SDOT_COLDESC:
 MOTOR VEHICLE STRUCK MOTOR VEHICLE, FRONT END AT ANGLE         83027
MOTOR VEHICLE STRUCK MOTOR VEHICLE, REAR END                   52488
MOTOR VEHICLE STRUCK MOTOR VEHICLE, LEFT SIDE SIDESWIPE         9776
MOTOR VEHICLE RAN OFF ROAD - HIT FIXED OBJECT                   8699
MOTOR VEHCILE STRUCK PEDESTRIAN                                 6368
MOTOR VEHICLE STRUCK MOTOR VEHICLE, LEFT SIDE AT ANGLE          5614
MOTOR VEHICLE STRUCK OBJECT IN ROAD                             4581
NOT ENOUGH INFORMATION / NOT APPLICABLE                         3112
MOTOR VEHICLE STRUCK PEDALCYCLIST, FRONT END AT ANGLE           3030
MOTOR VEHICLE STRUCK MOTOR VEHICLE, RIGHT SIDE SIDESWIPE        1567
MOTOR VEHICLE STRUCK MOTOR VEHICLE, RIGHT SIDE AT ANGLE         1370
```

```
PEDALCYCLIST STRUCK MOTOR VEHICLE FRONT END AT ANGLE            1292
MOTOR VEHICLE OVERTURNED IN ROAD                                 472
MOTOR VEHICLE STRUCK PEDALCYCLIST, REAR END                      180
PEDALCYCLIST STRUCK MOTOR VEHICLE LEFT SIDE SIDESWIPE            177
MOTOR VEHICLE RAN OFF ROAD - NO COLLISION                        160
PEDALCYCLIST STRUCK MOTOR VEHICLE REAR END                       134
MOTOR VEHICLE STRUCK PEDALCYCLIST, LEFT SIDE SIDESWIPE           122
DRIVERLESS VEHICLE RAN OFF ROAD - HIT FIXED OBJECT               106
DRIVERLESS VEHICLE STRUCK MOTOR VEHICLE FRONT END AT ANGLE       103
MOTOR VEHICLE STRUCK TRAIN                                       101
DRIVERLESS VEHICLE STRUCK MOTOR VEHICLE REAR END                  92
PEDALCYCLIST STRUCK PEDESTRIAN                                    74
PEDALCYCLIST OVERTURNED IN ROAD                                   67
DRIVERLESS VEHICLE STRUCK MOTOR VEHICLE LEFT SIDE AT ANGLE        53
PEDALCYCLIST STRUCK MOTOR VEHICLE RIGHT SIDE SIDESWIPE            50
PEDALCYCLIST STRUCK OBJECT IN ROAD                                23
MOTOR VEHICLE STRUCK PEDALCYCLIST, RIGHT SIDE SIDESWIPE           16
DRIVERLESS VEHICLE STRUCK MOTOR VEHICLE RIGHT SIDE AT ANGLE       12
PEDALCYCLIST STRUCK MOTOR VEHICLE LEFT SIDE AT ANGLE               9
DRIVERLESS VEHICLE STRUCK PEDESTRIAN                               8
PEDALCYCLIST STRUCK PEDALCYCLIST REAR END                          7
DRIVERLESS VEHICLE STRUCK MOTOR VEHICLE RIGHT SIDE SIDESWIPE       6
PEDALCYCLIST STRUCK PEDALCYCLIST FRONT END AT ANGLE                4
PEDALCYCLIST RAN OFF ROAD - HIT FIXED OBJECT                       4
DRIVERLESS VEHICLE STRUCK MOTOR VEHICLE LEFT SIDE SIDESWIPE        4
DRIVERLESS VEHICLE STRUCK OBJECT IN ROADWAY                        3
PEDALCYCLIST STRUCK MOTOR VEHICLE RIGHT SIDE AT ANGLE              2
DRIVERLESS VEHICLE RAN OFF ROAD - NO COLLISION                     1
Name: SDOT_COLDESC, dtype: int64

 ROADCOND:
 Dry               122159
Wet                46720
Unknown            11521
Ice                 1178
Snow/Slush           978
Other                123
Standing Water       108
Sand/Mud/Dirt         67
Oil                   60
Name: ROADCOND, dtype: int64
```

0.2 Exploring the Data

Now we have cleaned the data let's create some visuals to see what we are dealing with

[19]:

```
#This will look at the effects of the junction type
main_df.groupby(['JUNCTIONTYPE', 'SEVERITYDESC']).agg('size').unstack().
 ↪plot(kind = 'bar', legend=True, figsize=(11, 6), fontsize=11,␣
 ↪edgecolor='black',color=['blue', 'red',])
plt.ylim((-1000,80000))
```

[19]: (-1000.0, 80000.0)

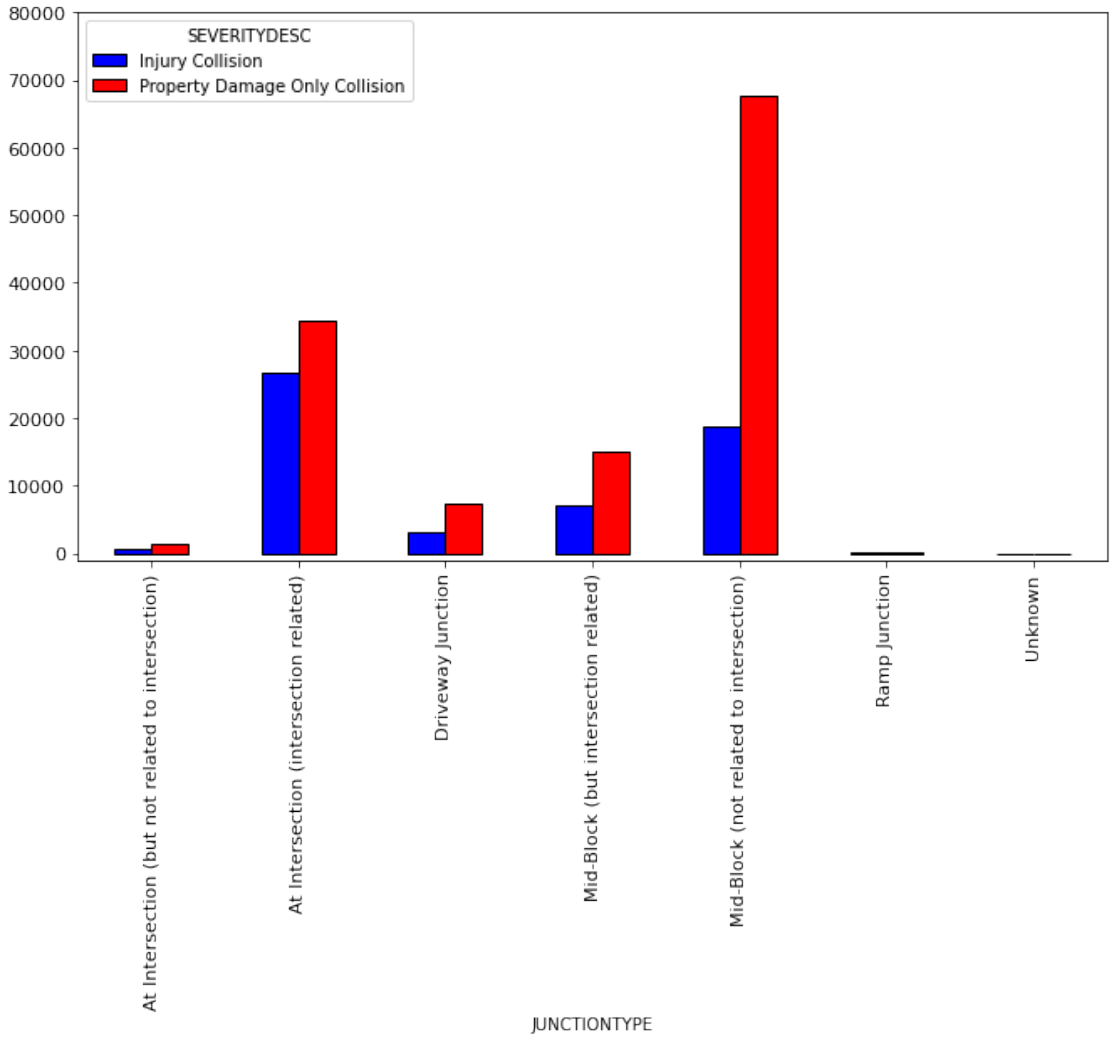

[20]:
```
#This looks at the effects of weather
main_df.groupby(['WEATHER', 'SEVERITYDESC']).agg('size').unstack().plot(kind =␣
 ↪'bar', figsize=(9,6), legend=True, fontsize=11, color=['green', 'blue'])
plt.ylim((-1000, 100000))
```

[20]: (-1000.0, 100000.0)

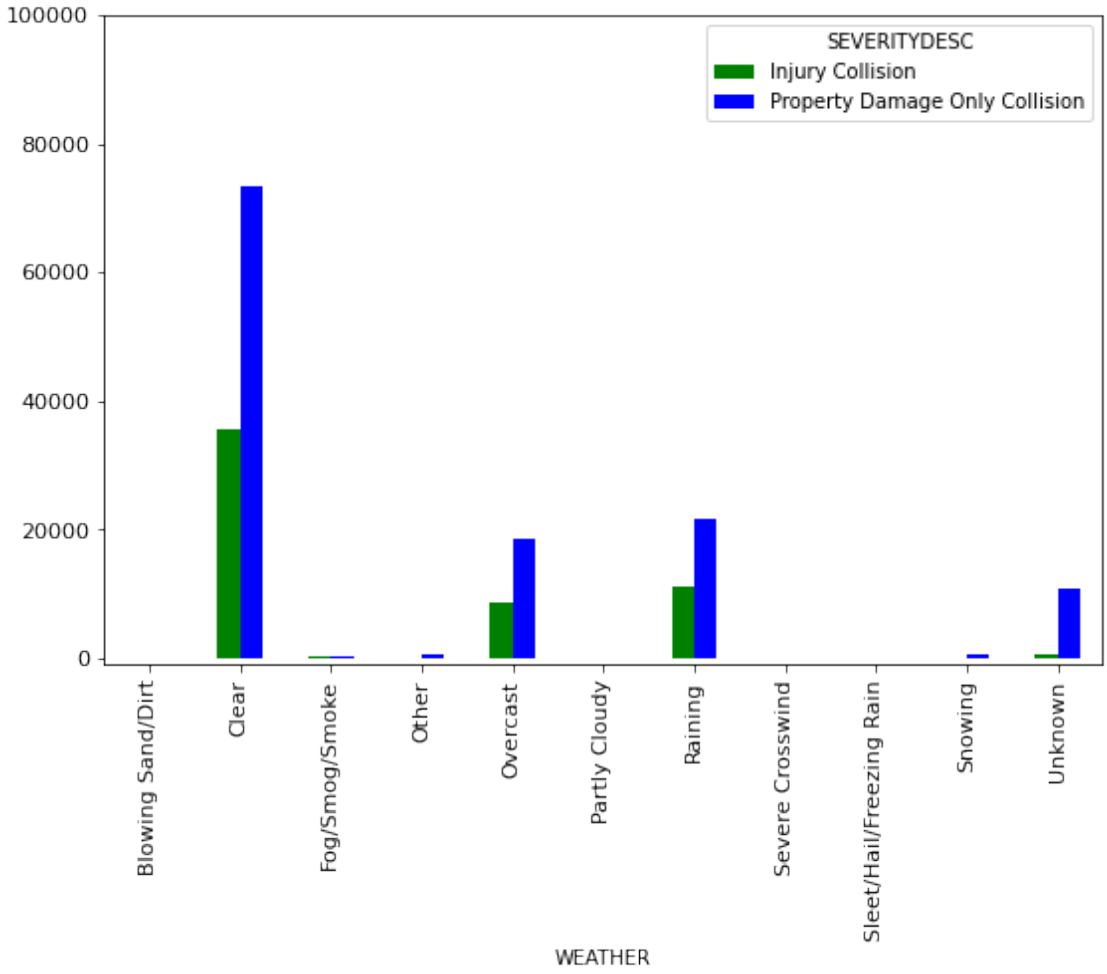

[21]:
```
X = np.arange(-5.0, 5.0, 0.1)
Y = 1.0 / (1.0 + np.exp(-X))

plt.plot(X,Y)
plt.ylabel('WEATHER')
plt.xlabel('SEVERITYCODE')
plt.show()
```

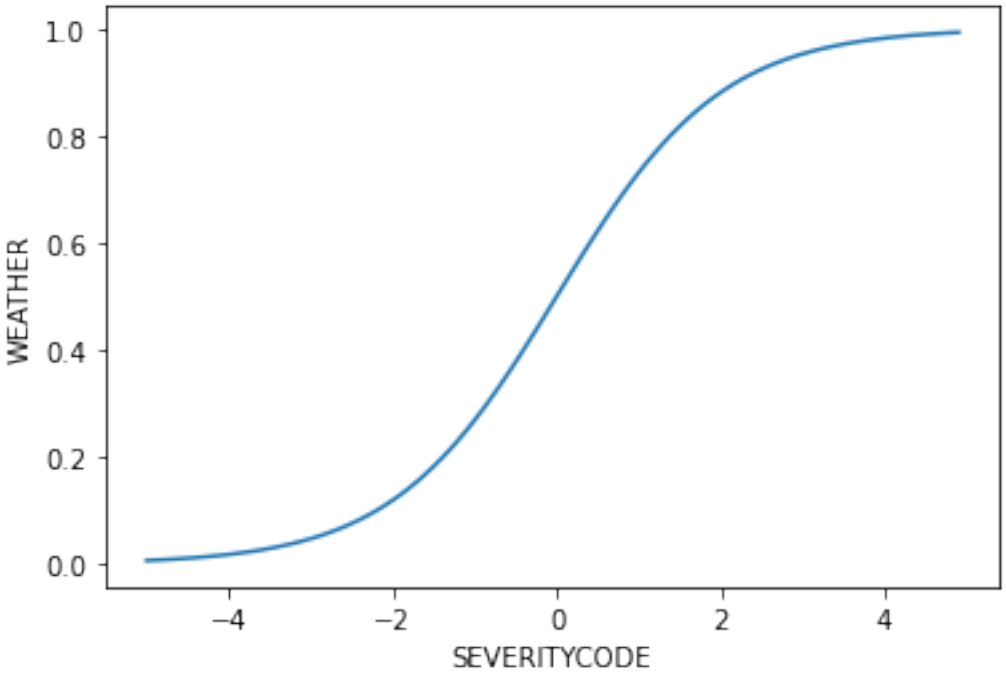

[22]:
```
#This looks at the effect of light condition
main_df.groupby(['LIGHTCOND', 'SEVERITYDESC']).agg('size').unstack().plot(kind␣
↪= 'bar', figsize=(8,6), legend=True, fontsize=12, color=['orange', 'blue'])

plt.ylim((-1000, 100000))
```

[22]: (-1000.0, 100000.0)

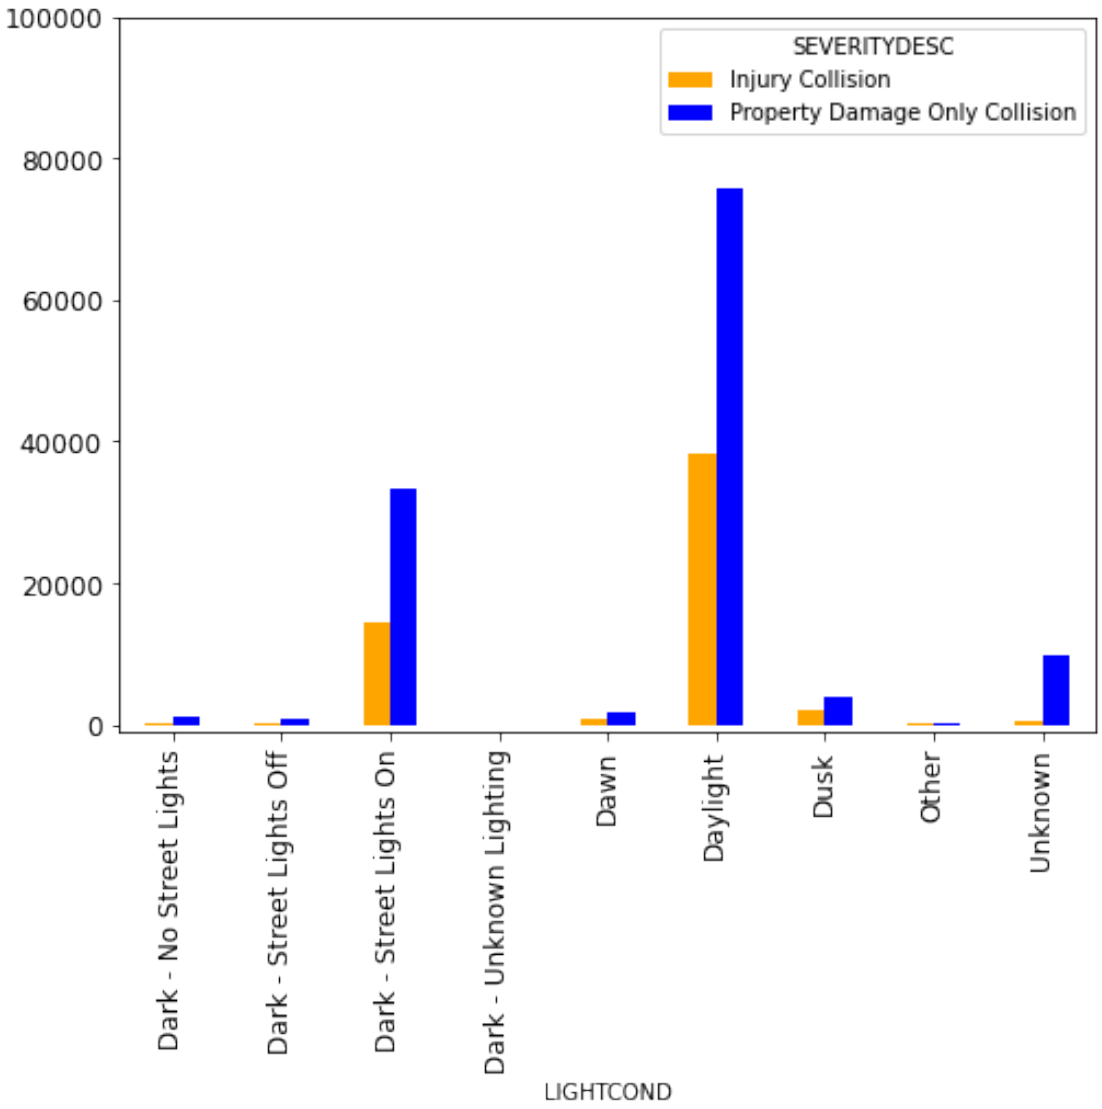

[23]:
```
#This looks at the effect of road conditions
sns.countplot(x = "ROADCOND" , data = main_df, )
plt.title("Effects of Roadcondition")
plt.xticks(rotation='vertical')
plt.show()
```

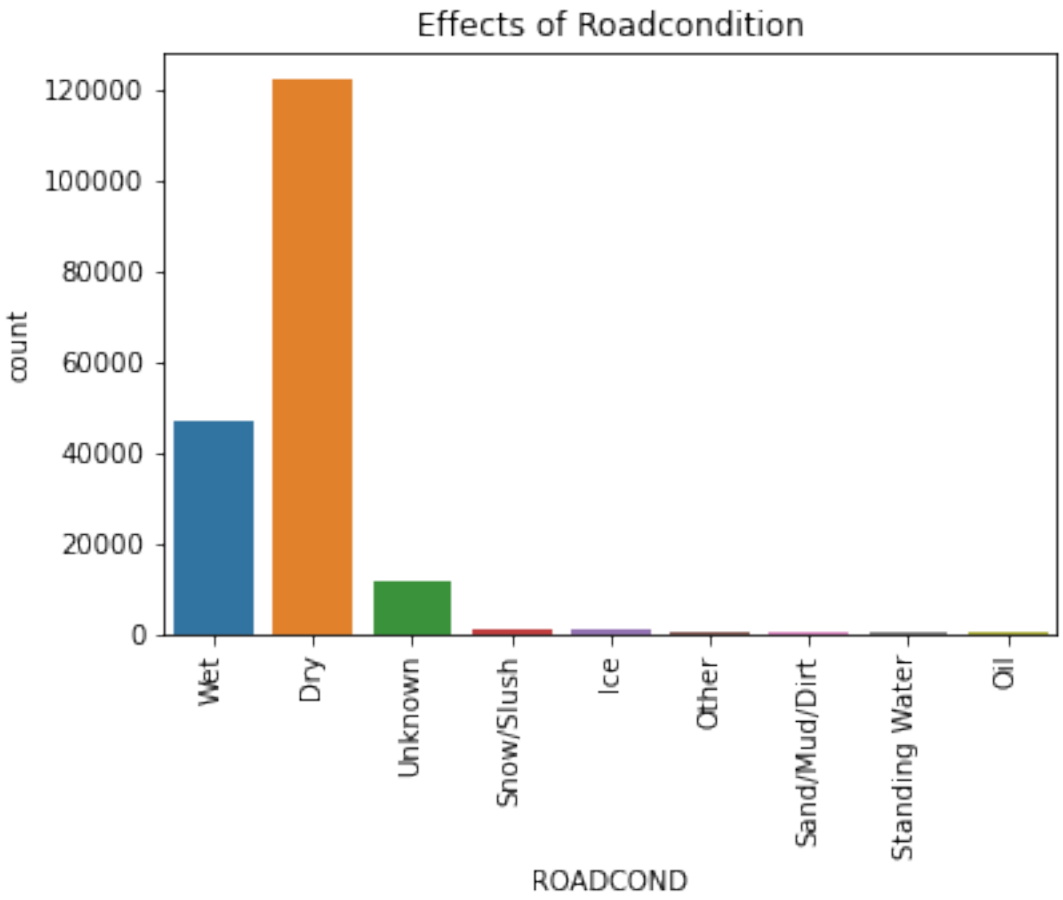

[24]:
```
#A breakdown in roadconditions numerically
main_df ['ROADCOND'].value_counts()
```

[24]:
```
Dry              122159
Wet               46720
Unknown           11521
Ice                1178
Snow/Slush          978
Other               123
Standing Water      108
Sand/Mud/Dirt        67
Oil                  60
Name: ROADCOND, dtype: int64
```

[26]:
```
from sklearn.model_selection import train_test_split

train, test = train_test_split(main_df, test_size=0.2)
```

[ ]: